

Modelos de Deep Learning

Entrenamiento de un MLP: GD y SGD

Universidad ORT Uruguay

1 de Setiembre, 2025

Funciones de Pérdida

- Regresión (MSE):

$$L(\hat{y}, y) = (\hat{y} - y)^2$$

- Clasificación binaria (BCE): la última activación es **sigmoidea**

$$L(\hat{y}, y) = - \left(y \log \hat{y} + (1 - y) \log(1 - \hat{y}) \right)$$

- Clasificación con K clases (CCE): la última activación es **softmax**

$$L(\hat{\mathbf{y}}, \mathbf{y}) = - \sum_{k=1}^O y_k \log \hat{y}_k$$

con $\mathbf{y} \in \{0, 1\}^O$ (one-hot encoding)

Costo Empírico

- Dada una función de pérdida $L(\hat{y}, y)$
- Para un (\mathbf{W}, \mathbf{b}) dado, la costo empírico es:

$$J(\mathbf{W}, \mathbf{b}) = \frac{1}{N} \sum_{i=1}^N L(f(\mathbf{x}_i; \mathbf{W}, \mathbf{b}), y_i) = \frac{1}{N} \sum_{i=1}^N L_i$$

donde $\hat{\mathbf{y}}_i = f(\mathbf{x}_i; \mathbf{W}, \mathbf{b})$ es la salida de la red para la entrada \mathbf{x}_i .

- Entrenar la red significa encontrar el argumento mínimo de J (ERM):

$$(\widehat{\mathbf{W}}, \widehat{\mathbf{b}}) = \arg \min_{(\mathbf{W}, \mathbf{b})} J(\mathbf{W}, \mathbf{b}).$$

Entrenamiento del MLP: Gradient Descent

- Problema de optimización - Algoritmo de **descenso por gradiente**:

$$(\mathbf{W}_0, \mathbf{b}_0) = \text{inicializar random}$$

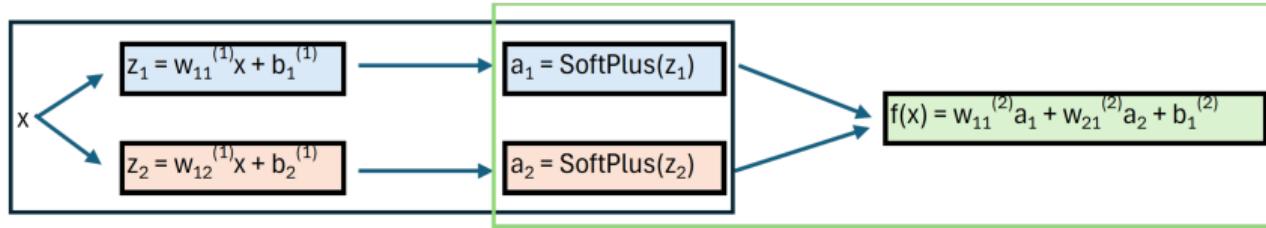
$$(\mathbf{W}_{k+1}, \mathbf{b}_{k+1}) = (\mathbf{W}_k, \mathbf{b}_k) - \eta \cdot \nabla J(\mathbf{W}_k, \mathbf{b}_k)$$

donde η es la **tasa de aprendizaje** (learning rate)

- Debemos **calcular** el **gradiente de J** en cada paso.
- **Linealidad:**

$$\nabla J(\mathbf{W}, \mathbf{b}) = \frac{1}{N} \sum_{i=1}^N \nabla L(f(\mathbf{x}_i; \mathbf{W}, \mathbf{b}), y_i) = \underset{(\mathbf{x}, y) \sim T}{\text{Promedio}} [\nabla L(f(\mathbf{x}; \mathbf{W}, \mathbf{b}), y)]$$

Ejemplo: MLP



■ La red es

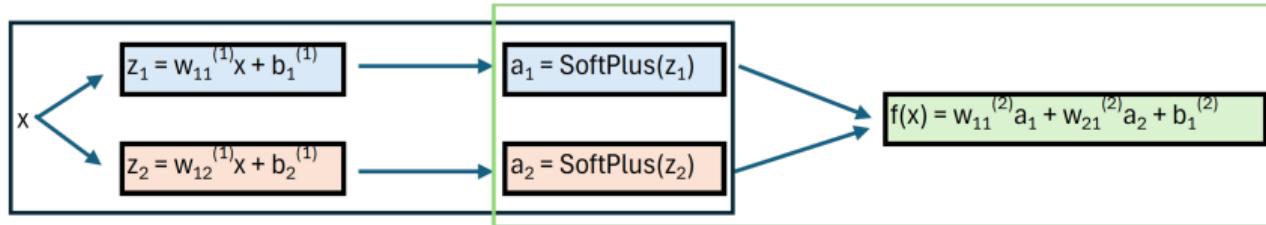
$$\hat{y} = \text{softplus}(x\mathbf{W}^{(1)} + \mathbf{b}^{(1)})\mathbf{W}^{(2)} + \mathbf{b}^{(2)}$$

■ $\mathbf{W}^{(1)} \sim (1, 2)$, $\mathbf{b}^{(1)} \sim (1, 2)$, $\mathbf{W}^{(2)} \sim (2, 1)$, $\mathbf{b}^{(2)} \sim (0)$

■ El costo es

$$\text{MSE} = \frac{1}{N} \sum_{i=1}^N (\hat{y}_i - y_i)^2$$

Ejemplo: cálculo de derivadas

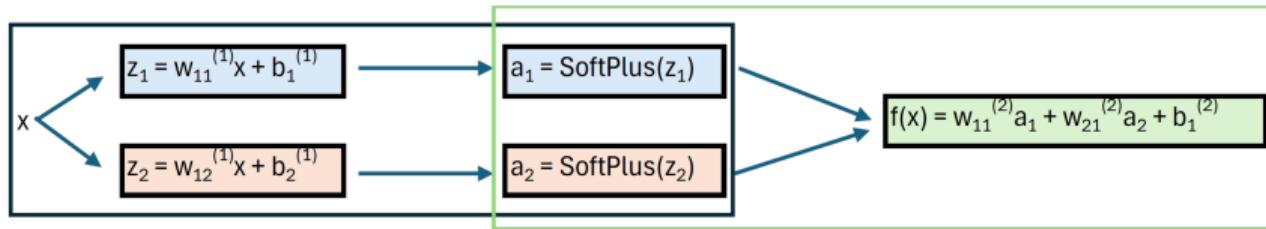


■ **Objetivo:** Calcular $\nabla L(f(x), y)$ para (x, y) fijo.

■ $\frac{\partial L}{\partial \hat{y}} = 2(\hat{y} - y)$

■ $\frac{\partial \text{softplus}(z)}{\partial z} = \sigma(z)$ (sigmoidea)

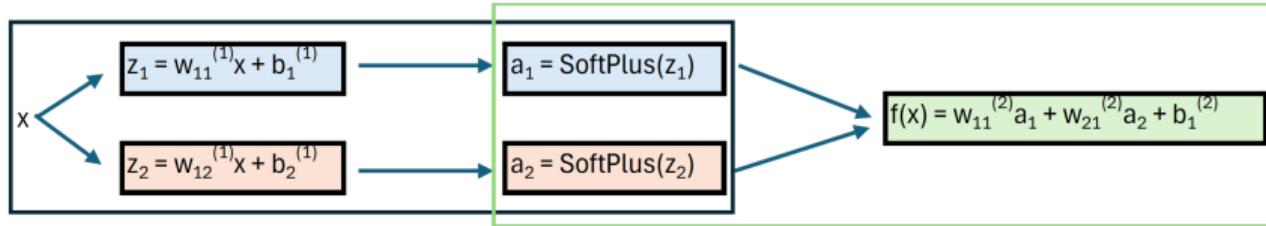
Ejemplo: cálculo de derivadas



■ $\frac{\partial \hat{y}}{\partial b_1^{(2)}} = 1$

■ $\frac{\partial L}{\partial b_1^{(2)}} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial b_1^{(2)}} \implies \frac{\partial L}{\partial b_1^{(2)}} = 2(\hat{y} - y)$

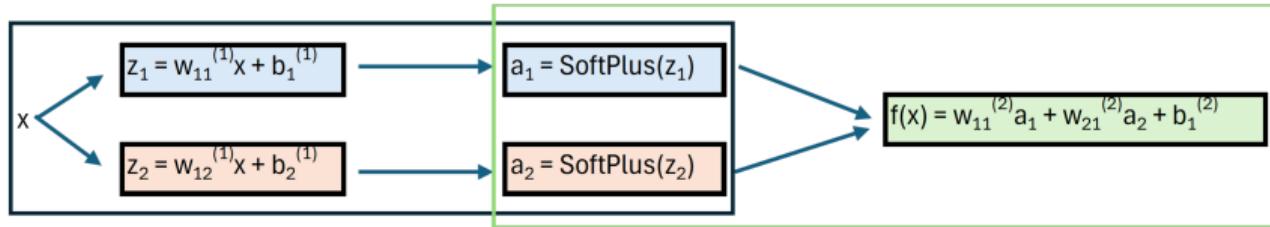
Ejemplo: cálculo de derivadas



- $\frac{\partial L}{\partial w_{11}^{(2)}} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial w_{11}^{(2)}} \implies \frac{\partial L}{\partial w_{11}^{(2)}} = 2(\hat{y} - y)a_1.$
- $\frac{\partial L}{\partial w_{21}^{(2)}} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial w_{21}^{(2)}} \implies \frac{\partial L}{\partial w_{21}^{(2)}} = 2(\hat{y} - y)a_2.$
- El Jacobiano es:

$$\frac{\partial L}{\partial \mathbf{W}^{(2)}} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial \mathbf{W}^{(2)}} = 2(\hat{y} - y) [a_1 \quad a_2] = 2(\hat{y} - y)\mathbf{a}^\top$$

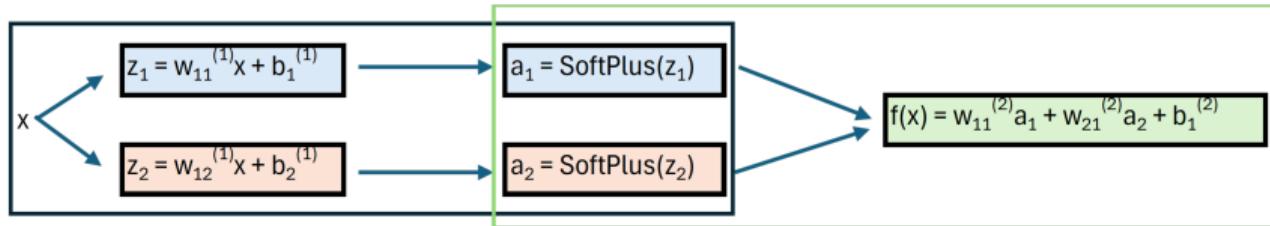
Ejemplo: cálculo de derivadas



- $\frac{\partial L}{\partial a_1} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial a_1} \implies \frac{\partial L}{\partial a_1} = 2(\hat{y} - y)w_{11}^{(2)}$.
- $\frac{\partial L}{\partial a_2} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial a_2} \implies \frac{\partial L}{\partial a_2} = 2(\hat{y} - y)w_{21}^{(2)}$.
- El Jacobiano es:

$$\frac{\partial L}{\partial \mathbf{a}} = 2(\hat{y} - y) \begin{bmatrix} w_{11}^{(2)} & w_{21}^{(2)} \end{bmatrix} = 2(\hat{y} - y) (\mathbf{W}^{(2)})^\top$$

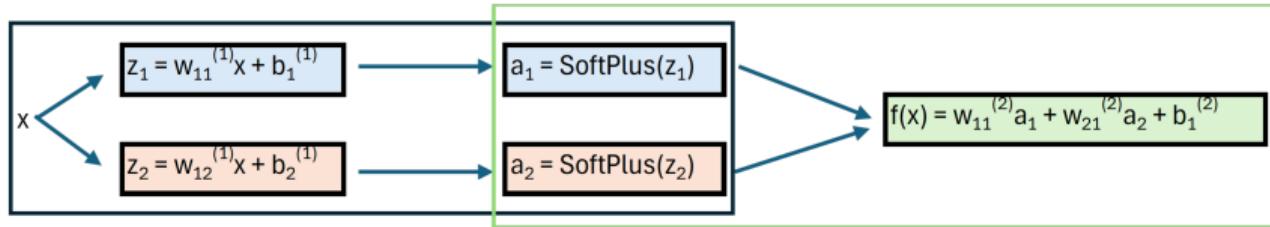
Ejemplo: cálculo de derivadas



- $\frac{\partial L}{\partial z_1} = \frac{\partial L}{\partial a_1} \frac{\partial a_1}{\partial z_1} \implies \frac{\partial L}{\partial z_1} = 2(\hat{y} - y)w_{11}^{(2)}\sigma(z_1).$
- $\frac{\partial L}{\partial z_2} = \frac{\partial L}{\partial a_2} \frac{\partial a_2}{\partial z_2} \implies \frac{\partial L}{\partial z_2} = 2(\hat{y} - y)w_{21}^{(2)}\sigma(z_2).$
- El Jacobiano es:

$$\begin{aligned}\frac{\partial L}{\partial \mathbf{z}} &= \frac{\partial L}{\partial \mathbf{a}} \frac{\partial \mathbf{a}}{\partial \mathbf{z}} = 2(\hat{y} - y) \begin{bmatrix} w_{11}^{(2)} & w_{21}^{(2)} \end{bmatrix} \begin{bmatrix} \sigma(z_1) & 0 \\ 0 & \sigma(z_2) \end{bmatrix} \\ &= 2(\hat{y} - y) (\mathbf{W}^{(2)})^\top \sigma(\mathbf{z})\end{aligned}$$

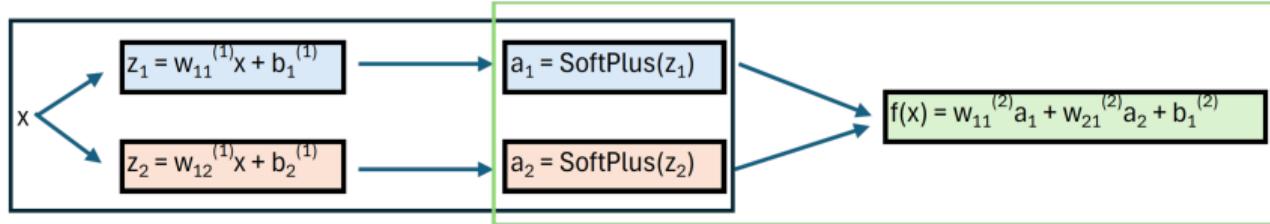
Ejemplo: cálculo de derivadas



- $\frac{\partial L}{\partial b_1^{(1)}} = \frac{\partial L}{\partial z_1} \frac{\partial z_1}{\partial b_1^{(1)}} \implies \frac{\partial L}{\partial b_1^{(1)}} = 2(\hat{y} - y)w_{11}^{(2)}\sigma(z_1).$
- $\frac{\partial L}{\partial b_2^{(1)}} = \frac{\partial L}{\partial z_2} \frac{\partial z_2}{\partial b_2^{(1)}} \implies \frac{\partial L}{\partial b_2^{(1)}} = 2(\hat{y} - y)w_{21}^{(2)}\sigma(z_2).$
- El Jacobiano es:

$$\frac{\partial L}{\partial \mathbf{b}^{(1)}} = \frac{\partial L}{\partial \mathbf{z}} \underbrace{\frac{\partial \mathbf{z}}{\partial \mathbf{b}^{(1)}}}_{\text{Id}} = 2(\hat{y} - y) (\mathbf{W}^{(2)})^\top \sigma(\mathbf{z})$$

Ejemplo: cálculo de derivadas



- $\frac{\partial L}{\partial w_{11}^{(1)}} = \frac{\partial L}{\partial z_1} \frac{\partial z_1}{\partial w_{11}^{(1)}} \implies \frac{\partial L}{\partial w_{11}^{(1)}} = 2(\hat{y} - y)w_{11}^{(2)}\sigma(z_1)x.$
- $\frac{\partial L}{\partial w_{12}^{(1)}} = \frac{\partial L}{\partial z_2} \frac{\partial z_2}{\partial w_{12}^{(1)}} \implies \frac{\partial L}{\partial w_{12}^{(1)}} = 2(\hat{y} - y)w_{12}^{(2)}\sigma(z_2)x.$
- El Jacobiano es:

$$\begin{aligned}\frac{\partial L}{\partial \mathbf{W}^{(1)}} &= \frac{\partial L}{\partial \mathbf{z}} \frac{\partial \mathbf{z}}{\partial \mathbf{W}^{(1)}} = 2(\hat{y} - y) \begin{bmatrix} w_{11}^{(2)}\sigma(z_1) & w_{21}^{(2)}\sigma(z_2) \end{bmatrix} \begin{bmatrix} x & 0 \\ 0 & x \end{bmatrix} \\ &= 2(\hat{y} - y) \cdot \mathbf{W}^{(2)}\sigma(\mathbf{z}) \cdot x\end{aligned}$$

Ejemplo: Resumiendo el cálculo directo

- $\frac{\partial L}{\partial b_1^{(2)}} = 2(\hat{y} - y)$
- $\frac{\partial L}{\partial \mathbf{W}^{(2)}} = 2(\hat{y} - y)\mathbf{a}$
- $\frac{\partial L}{\partial \mathbf{a}} = 2(\hat{y} - y) (\mathbf{W}^{(2)})^\top$
- $\frac{\partial L}{\partial \mathbf{z}} = 2(\hat{y} - y) (\mathbf{W}^{(2)})^\top \sigma(\mathbf{z})$
- $\frac{\partial L}{\partial \mathbf{b}^{(1)}} = 2(\hat{y} - y) (\mathbf{W}^{(2)})^\top \sigma(\mathbf{z})$
- $\frac{\partial L}{\partial \mathbf{W}^{(1)}} = 2(\hat{y} - y) \cdot (\mathbf{W}^{(2)})^\top \sigma(\mathbf{z}) \cdot x$

Ejemplo: operando en batch (útil para la notebook)

- $\frac{\partial J}{\partial \hat{\mathbf{y}}} = \frac{2}{N}(\hat{\mathbf{y}} - \mathbf{y})^\top \sim (1, N)$
- $\frac{\partial J}{\partial \mathbf{b}^{(2)}} = \frac{2}{N}(\hat{\mathbf{y}} - \mathbf{y})^\top \mathbf{1}_N \sim (0)$
- $\frac{\partial J}{\partial \mathbf{W}^{(2)}} = \frac{2}{N}(\hat{\mathbf{y}} - \mathbf{y})^\top \mathbf{A} \sim (1, 2)$ donde $\mathbf{A} = \begin{bmatrix} \mathbf{a}_1^\top \\ \vdots \\ \mathbf{a}_N^\top \end{bmatrix}$ (!) al implementar trasponer
- $\frac{\partial J}{\partial \mathbf{b}^{(1)}} = \frac{2}{N} \mathbf{1}_N^\top \left[\left[(\hat{\mathbf{y}} - \mathbf{y}) (\mathbf{W}^{(2)})^\top \right] \odot \sigma(\mathbf{Z}) \right] \sim (1, 2)$ con $\sigma(\mathbf{Z}) = \begin{bmatrix} \sigma(\mathbf{z}_1)^\top \\ \vdots \\ \sigma(\mathbf{z}_N)^\top \end{bmatrix}$
- $\frac{\partial J}{\partial \mathbf{W}^{(1)}} = \frac{2}{N} \mathbf{X}^\top \left[\left[(\hat{\mathbf{y}} - \mathbf{y}) (\mathbf{W}^{(2)})^\top \right] \odot \sigma(\mathbf{Z}) \right] \sim (1, 2)$ con $\mathbf{X} = \begin{bmatrix} x_1 \\ \vdots \\ x_N \end{bmatrix}$

Ejemplo: volviendo a ∇L pero en reversa

- $\delta_2 = \frac{\partial L}{\partial \hat{y}} = 2(\hat{y} - y)$
- $\delta_1 = \frac{\partial L}{\partial \mathbf{z}} = 2(\hat{y} - y) (\mathbf{W}^{(2)})^\top \sigma(\mathbf{z})$

De aquí se obtienen las derivadas de interés:

- En la capa 2:
 - $\frac{\partial L}{\partial \mathbf{b}^{(2)}} = \delta_2$
 - $\frac{\partial L}{\partial \mathbf{W}^{(2)}} = \delta_2 \mathbf{a}^\top$
- En la capa 1:
 - $\frac{\partial L}{\partial \mathbf{b}^{(1)}} = \delta_1$
 - $\frac{\partial L}{\partial \mathbf{W}^{(1)}} = \delta_1 \mathbf{x}$

Backpropagation para un MLP

Objetivo: Calcular $\nabla L(\hat{y}, y)$ eficientemente para un MLP con K capas.

- $\mathbf{z}_k = \mathbf{a}_{k-1}^\top \mathbf{W}_k + \mathbf{b}_k^\top$ (pre-activación)
- $\mathbf{a}_k = A(\mathbf{z}_k)$ (post-activación)
- $\mathbf{a}^0 = \mathbf{x}$

Forward pass:

- Para $k = 1$ hasta K : calcular \mathbf{z}_k y \mathbf{a}_k

Backward pass (errores):

- Calcular $\delta_K = A'(\mathbf{z}_K) \nabla_{\mathbf{a}_K} L$
- Para $k = K - 1$ hasta 1: $\delta_k = \mathbf{W}_{k+1}^\top A'(\mathbf{z}_k) \delta_{k+1}$

Backpropagation para un MLP

Gradientes:

Para cada capa $k = 1, \dots, K$:

$$\nabla_{\mathbf{W}_k} L = \delta_k \cdot \mathbf{a}_{k-1}^\top$$

$$\nabla_{\mathbf{b}_k} L = \delta_k$$

Linealidad:

$$\nabla J = \frac{1}{N} \sum_{i=1}^N \nabla L_i$$

Actualización:

$$\mathbf{W}_k \leftarrow \mathbf{W}_k - \eta \nabla_{\mathbf{W}_k} J, \quad \mathbf{b}_k \leftarrow \mathbf{b}_k - \eta \nabla_{\mathbf{b}_k} J$$

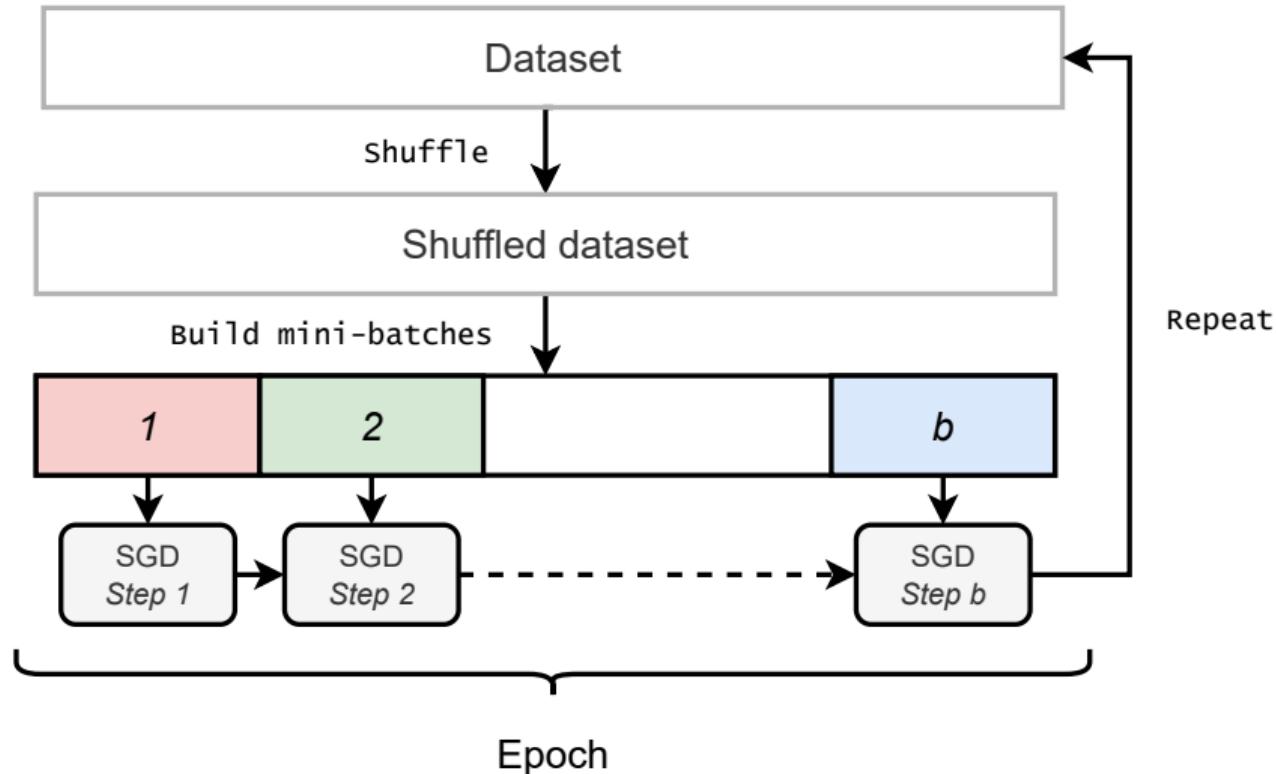
Entrenamiento del MLP: Stochastic Gradient Descent

Algoritmo básico: Descenso del gradiente estocástico (SGD)

- Comenzar con $(\mathbf{W}_0, \mathbf{b}_0)$ random
- En cada **step** t : seleccionar $(x, y) \in T$ aleatoriamente y actualizar
$$(\mathbf{W}_{t+1}, \mathbf{b}_{t+1}) = (\mathbf{W}_t, \mathbf{b}_t) - \eta \nabla L(f(\mathbf{x}; \mathbf{W}_t, \mathbf{b}_t), y) \quad \eta \text{ es el learning rate}$$

En la práctica se actualiza por **mini-batch** (se aprovecha el paralelismo)

SGD con mini-batches



Entrenamiento del MLP: SGD con mini-batches

- Comenzar con $(\mathbf{W}_0, \mathbf{b}_0)$ random
- En cada epoch:
 - Seleccionar el batch B
 - Actualizar pesos:

$$(\mathbf{W}, \mathbf{b}) \leftarrow (\mathbf{W}, \mathbf{b}) - \eta \nabla J_B(f(\mathbf{x}; \mathbf{W}, \mathbf{b}), y) \quad \eta \text{ es el learning rate}$$

donde

$$J_B = \frac{1}{n} \sum_{(x_i, y_i) \in B} L_i$$